

```
In [1]: # Importing all the required packages
import warnings
warnings.filterwarnings("ignore")
from sklearn.datasets import load_boston
from random import seed
from random import randrange
from csv import reader
from math import sqrt
from sklearn import preprocessing
import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
from prettytable import PrettyTable
from sklearn.linear_model import SGDRegressor
from sklearn import preprocessing
from sklearn.metrics import mean_squared_error
```

Reading the dataset

```
In [2]: X = load_boston().data
        Y = load_boston().target
```

```
In [3]: #Description of dataset
print(load_boston().DESCR)
```

```
Boston House Prices dataset
=====
```

```
Notes
```

```
-----
```

```
Data Set Characteristics:
```

```
:Number of Instances: 506
```

```
:Number of Attributes: 13 numeric/categorical predictive
```

```
:Median Value (attribute 14) is usually the target
```

```
:Attribute Information (in order):
```

```
  - CRIM      per capita crime rate by town
  - ZN        proportion of residential land zoned for lots over 25,000 sq.ft.
  - INDUS     proportion of non-retail business acres per town
  - CHAS      Charles River dummy variable (= 1 if tract bounds river; 0 otherwise)
  - NOX       nitric oxides concentration (parts per 10 million)
  - RM        average number of rooms per dwelling
  - AGE       proportion of owner-occupied units built prior to 1940
  - DIS       weighted distances to five Boston employment centres
  - RAD       index of accessibility to radial highways
  - TAX       full-value property-tax rate per $10,000
  - PTRATIO   pupil-teacher ratio by town
  - B         1000(Bk - 0.63)^2 where Bk is the proportion of blacks by town
  - LSTAT     % lower status of the population
  - MEDV      Median value of owner-occupied homes in $1000's
```

```
:Missing Attribute Values: None
```

```
:Creator: Harrison, D. and Rubinfeld, D.L.
```

```
This is a copy of UCI ML housing dataset.
```

```
http://archive.ics.uci.edu/ml/datasets/Housing (http://archive.ics.uci.edu/ml/datasets/Housing)
```

This dataset was taken from the StatLib library which is maintained at Carnegie Mellon University.

The Boston house-price data of Harrison, D. and Rubinfeld, D.L. 'Hedonic prices and the demand for clean air', J. Environ. Economics & Management, vol.5, 81-102, 1978. Used in Belsley, Kuh & Welsch, 'Regression diagnostics ...', Wiley, 1980. N.B. Various transformations are used in the table on pages 244-261 of the latter.

The Boston house-price data has been used in many machine learning papers that address regression problems.

****References****

- Belsley, Kuh & Welsch, 'Regression diagnostics: Identifying Influential Data and Sources of Collinearity', Wiley, 1980. 244-261.
- Quinlan, R. (1993). Combining Instance-Based and Model-Based Learning. In Proceedings on the Tenth International Conference of Machine Learning, 236-243, University of Massachusetts, Amherst. Morgan Kaufmann.
- many more! (see <http://archive.ics.uci.edu/ml/datasets/Housing>) (<http://archive.ics.uci.edu/ml/datasets/Housing>)

Preprocessing the dataset

```
In [4]: # As the pre-loaded dataset is a numpy array, we are converting it to dataframe
# https://stackoverflow.com/questions/20763012/creating-a-pandas-dataframe-from-numpy-array
bos_df = pd.DataFrame(data=load_boston().data, columns=load_boston().feature_names)
print("The shape of data frame is", bos_df.shape)
print("\nThe top 5 rows of the dataframe are")
bos_df.head()
```

The shape of data frame is (506, 13)

The top 5 rows of the dataframe are

Out[4]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [5]: # Inserting the target variable in the dataframe
bos_df['price'] = load_boston().target
bos_df.head()
```

Out[5]:

	CRIM	ZN	INDUS	CHAS	NOX	RM	AGE	DIS	RAD	TAX	PTRATIO	B	LSTAT
0	0.00632	18.0	2.31	0.0	0.538	6.575	65.2	4.0900	1.0	296.0	15.3	396.90	4.98
1	0.02731	0.0	7.07	0.0	0.469	6.421	78.9	4.9671	2.0	242.0	17.8	396.90	9.14
2	0.02729	0.0	7.07	0.0	0.469	7.185	61.1	4.9671	2.0	242.0	17.8	392.83	4.03
3	0.03237	0.0	2.18	0.0	0.458	6.998	45.8	6.0622	3.0	222.0	18.7	394.63	2.94
4	0.06905	0.0	2.18	0.0	0.458	7.147	54.2	6.0622	3.0	222.0	18.7	396.90	5.33

```
In [6]: bos_df.columns
```

```
Out[6]: Index(['CRIM', 'ZN', 'INDUS', 'CHAS', 'NOX', 'RM', 'AGE', 'DIS', 'RAD', 'TAX',  
             'PTRATIO', 'B', 'LSTAT', 'price'],  
            dtype='object')
```

```
In [7]: # Checking if any null value is present in dataset  
bos_df.isnull().sum()
```

```
Out[7]: CRIM      0  
        ZN        0  
        INDUS    0  
        CHAS     0  
        NOX      0  
        RM       0  
        AGE      0  
        DIS      0  
        RAD      0  
        TAX      0  
        PTRATIO  0  
        B        0  
        LSTAT    0  
        price    0  
        dtype: int64
```

Splitting the dataset into train and test

```
In [8]: from sklearn.model_selection import train_test_split  
X_train, X_test, y_train, y_test = train_test_split(X, Y, test_size=0.33, random_
```

```
In [9]: print("X_train shape is", X_train.shape)  
        print("y_train shape is", y_train.shape)  
        print("X_test shape is", X_test.shape)  
        print("y_test shape is", y_test.shape)
```

```
X_train shape is (339, 13)  
y_train shape is (339,)  
X_test shape is (167, 13)  
y_test shape is (167,)
```

Standardizing the dataset

```
In [10]: from sklearn.preprocessing import StandardScaler  
S = StandardScaler()  
X_train = S.fit_transform(X_train)  
X_test = S.transform(X_test)
```

References used to implement SGD on Linear regression

1. <https://machinelearningmastery.com/implement-linear-regression-stochastic-gradient-descent-scratch-python/> (<https://machinelearningmastery.com/implement-linear-regression-stochastic-gradient-descent-scratch-python/>)
2. <https://towardsdatascience.com/step-by-step-tutorial-on-linear-regression-with-stochastic-gradient-descent-1d35b088a843> (<https://towardsdatascience.com/step-by-step-tutorial-on-linear-regression-with-stochastic-gradient-descent-1d35b088a843>)

Implementing SGD from scratch

```
In [11]: def pred(X,w,b):  
  
    y_hat = []  
    for i in range(0,len(X)):  
        y = np.dot(X[i],w)+b  
        #https://www.geeksforgeeks.org/numpy-astype-in-python  
        y = np.asscalar(y)  
        y_hat.append(y)  
    y_hat = np.array(y_hat)  
  
    return y_hat
```

```

In [12]: def SGD(X, y, l_rate, epochs, l_rate_var):

    w=np.random.randn(X.shape[1],1)
    b=np.random.randn(1,1)

    n=X.shape[0]
    for epoch in range(1,epochs+1):
        sum_error=0

        for i in range(n):
            batch=np.random.randint(0,n)
            x_batch=X[batch,:].reshape(1,X.shape[1])
            y_batch=y[batch].reshape(1,1)

            y_pred=np.dot(x_batch,w)+b
            error=y_pred-y_batch

            sum_error += error**2
            dw=x_batch.T.dot((y_pred-y_batch))# Arrived logically by definition
            db=(y_pred-y_batch)

            w=w-(2/n)*l_rate*(dw)
            b=b-(2/n)*l_rate*(db)
        sum_error=sum_error/n
        print("Epoch :{0}  Total_error:{1}  lr_rate:{2}".format(epoch, np.round(sum_error,2), l_rate))

        if l_rate_var == 'constant':
            pass
        else :
            l_rate = l_rate/2

    return w,b

```

Executing Sgd with constant learning rate for all epochs


```
In [13]: w_best,b_best = SGD(X_train,y_train,l_rate=0.01, epochs=500, l_rate_var = 'constant')
Best_weights = pd.DataFrame(w_best,columns=['SGD_constant_l_rate'])
```

```
Epoch :96 Total_error:[[36.83]] lr_rate:0.01
Epoch :97 Total_error:[[41.87]] lr_rate:0.01
Epoch :98 Total_error:[[32.98]] lr_rate:0.01
Epoch :99 Total_error:[[43.62]] lr_rate:0.01
Epoch :100 Total_error:[[29.15]] lr_rate:0.01
Epoch :101 Total_error:[[42.62]] lr_rate:0.01
Epoch :102 Total_error:[[32.85]] lr_rate:0.01
Epoch :103 Total_error:[[35.53]] lr_rate:0.01
Epoch :104 Total_error:[[32.04]] lr_rate:0.01
Epoch :105 Total_error:[[26.89]] lr_rate:0.01
Epoch :106 Total_error:[[36.21]] lr_rate:0.01
Epoch :107 Total_error:[[34.43]] lr_rate:0.01
Epoch :108 Total_error:[[27.79]] lr_rate:0.01
Epoch :109 Total_error:[[32.18]] lr_rate:0.01
Epoch :110 Total_error:[[32.78]] lr_rate:0.01

Epoch :111 Total_error:[[39.03]] lr_rate:0.01
Epoch :112 Total_error:[[34.41]] lr_rate:0.01
Epoch :113 Total_error:[[32.34]] lr_rate:0.01
Epoch :114 Total_error:[[27.77]] lr_rate:0.01
```

```
In [14]: # Plotting line plot for actual vs model predicted
# https://stackoverflow.com/questions/31069191/simple-line-plots-using-seaborn
import seaborn as sns
y_hat_1 = pred(X_test, w_best, b_best)
plt.scatter(x=y_test,y=y_hat_1)
plt.xlabel('Actual value')
plt.ylabel('predicted value')
ms_er1 = mean_squared_error(y_test,y_hat_1)
print("Mean squared error",ms_er1)
```

Mean squared error 21.186219541861036



Observations

1. From above we observe that loss is gradually reducing if learning rate is kept constant

Executing SGD with different learning rate for all epochs

```
In [15]: w_best,b_best=SGD(X_train,y_train,l_rate=1, epochs=500, l_rate_var = 'not constant')
Best_weights['SGD_varied_l_rate'] = pd.DataFrame(w_best,columns=['SGD_varied_l_rate'])
```

```
Epoch :105 Total_error:[[18.89]] lr_rate:4.930380657631324e-32
Epoch :106 Total_error:[[28.16]] lr_rate:2.465190328815662e-32


Epoch :107 Total_error:[[23.7]] lr_rate:1.232595164407831e-32
Epoch :108 Total_error:[[22.73]] lr_rate:6.162975822039155e-33
Epoch :109 Total_error:[[22.93]] lr_rate:3.0814879110195774e-33
Epoch :110 Total_error:[[22.31]] lr_rate:1.5407439555097887e-33
Epoch :111 Total_error:[[28.87]] lr_rate:7.703719777548943e-34
Epoch :112 Total_error:[[25.11]] lr_rate:3.85185988774472e-34
Epoch :113 Total_error:[[28.09]] lr_rate:1.925929944387236e-34
Epoch :114 Total_error:[[28.15]] lr_rate:9.62964972193618e-35
Epoch :115 Total_error:[[29.96]] lr_rate:4.81482486096809e-35
Epoch :116 Total_error:[[18.43]] lr_rate:2.407412430484045e-35
Epoch :117 Total_error:[[25.61]] lr_rate:1.2037062152420224e-35
Epoch :118 Total_error:[[21.77]] lr_rate:6.018531076210112e-36
Epoch :119 Total_error:[[29.13]] lr_rate:3.009265538105056e-36
Epoch :120 Total_error:[[25.24]] lr_rate:1.504632769052528e-36
Epoch :121 Total_error:[[27.63]] lr_rate:7.52316384526264e-37
Epoch :122 Total_error:[[27.18]] lr_rate:3.76158192263132e-37
Epoch :123 Total_error:[[20.97]] lr_rate:1.88079096131566e-37
```

Observation

1. We notice that there is a drastic change from 2nd epoch itself compared to the first approach i.e with constant learning rate


```
In [16]: # Plotting line plot for actual vs model predicted
# https://stackoverflow.com/questions/31069191/simple-line-plots-using-seaborn
import seaborn as sns
y_hat_1 = pred(X_test, w_best, b_best)
plt.scatter(x=y_test,y=y_hat_1)
plt.xlabel('Actual value')
plt.ylabel('predicted value')
ms_er2 = mean_squared_error(y_test,y_hat_1)
print("Mean squared error",ms_er2)
```

Mean squared error 22.060186907739414



Executing sklearn's SGDRegressor


```
In [17]: from sklearn.linear_model import SGDRegressor
SGDR = SGDRegressor(loss='squared_loss', alpha=0.01, max_iter=500)
SGDR.fit(X_train,y_train)
y_hat = SGDR.predict(X_test)
Best_weights['SGDRegressor_scikit'] = SGDR.coef_
```

```
In [18]: print(mean_squared_error(y_test, SGDR.predict(X_test)))
```

20.85601918948433

```
In [19]: # Plotting line plot for actual vs model predicted
# https://stackoverflow.com/questions/31069191/simple-line-plots-using-seaborn
import seaborn as sns
y_hat_1 = SGDR.predict(X_test)
plt.scatter(x=y_test,y=y_hat_1)
plt.xlabel('Actual value')
plt.ylabel('predicted value')
ms_er3 = mean_squared_error(y_test,y_hat_1)
print("Mean squared error",ms_er3)
```

Mean squared error 20.85601918948433



```
In [20]: Best_weights
```

```
Out[20]:
```

	SGD_constant_l_rate	SGD_varied_l_rate	SGDRegressor_scikit
0	-0.840674	-0.795003	-0.942057
1	0.712502	0.655565	0.804730
2	-0.009761	-0.590807	0.314651
3	0.834394	1.314008	0.884120
4	-1.189070	-0.159156	-1.786730
5	2.872937	3.100020	2.825662
6	-0.429867	-0.575831	-0.362631
7	-2.711371	-2.332190	-2.903074
8	1.397033	1.688015	1.780323
9	-0.851373	-1.718437	-1.152708
10	-1.914551	-1.116304	-2.048570
11	1.040986	0.919219	1.047731
12	-4.049084	-3.715732	-3.889523

```
In [21]: print("Mean score for SGD_constant_l_rate",ms_er1)
          print("Mean score for SGD_varied_l_rate",ms_er2)
          print("Mean score for SGDRegressor_scikit",ms_er3)
```

```
Mean score for SGD_constant_l_rate 21.186219541861036
Mean score for SGD_varied_l_rate 22.060186907739414
Mean score for SGDRegressor_scikit 20.85601918948433
```